

DNA Sequencing and Data Analysis

Prof Noam Shomron
Hadas Volkov

Lecture 6, May 16, 2024

DNA Sequencing and Data Analysis

Sequence Mapping and Alignment

Thursday 18:30 to 21:00

Hangar H2

nshomron@gmail.com

hadas.volkov@post.runi.ac.il

Why Do We Need Sequence Mapping?

Determine the origin of an unknown sequence

Find homologous sequences

Determine genomic position of a sequence


Identify genomic variants between samples (variant calling)

Determine the function of a sequence (annotation)

Two Stages of Sequence Mapping

1. **SEARCH** -

Roughly find the position of the query in the DB




2. **ALIGN** -

Find the exact pairwise alignment of the query and the DB sequences

G	T	G	T	A	C	A	-	G
G	T	A	T	A	C	A	A	G


Pairwise Alignment

Hemoglobin Homologous




© 2016 Healthwise


A. Amino Acids with Electrically Charged Side Chains




B. Amino Acids with Polar Uncharged Side Chains



C. Special Cases



D. Amino Acids with Hydrophobic Side Chains



Pairwise Alignment

Hemoglobin Homologous

```
# NCBI Reference Sequence: NP_000508.1 (human hemoglobin subunit A)
r1 = skbio.Protein("MVLSPADKTNVKAAWGKVGAGHAGEYGAELERMFLSFPTTKTYFPHFDLSHGSAQVKGHGKKVADAL\
TNAVAHVDDMPNALSALSDLHAHKLRVDPVNFKLLSHCLLVTLAAHLPAEFTPAVHASLDKFLASVSTVLTISKYR")

# NCBI Reference Sequence: NP_001004376.1 (chicken hemoglobin subunit A)
r2 = skbio.Protein("MVLSAADKNNVKGIFTKIAGHAEYGAETLERMFTTYPPTKTYFPHFDLSHGSAQIKGHGKKVVAAL\
IEAANHIDDIAGTLSKLSDLHAHKLRVDPVNFKLLGQCFLVVVAIHHPAALTPEVHASLDKFLCAVGTVLTAKYR")

# GenBank: QFF91579.1 (sei whale hemoglobin subunit A)
q1 = skbio.Protein("MVLFPADKSNVKATWAKIGNHGAEYGAELERMFMNFPSTKTYFPHFDLGHDSAQVKGHGKKVADAL\
TKAAGHMDNLLDALSDLSDLHAHKLRVDPVNFKLLSHCLLVTLALHLPAEFTPSVHASLDKFLASVSTVLTISKYR")
```

Pairwise Alignment

The Hamming Distance

Hamming distance is the number of symbols or positions of two strings at which their corresponding characters are different

```
def hamming_distance(string1, string2):  
    if (len(string1) != len(string2)):  
        raise Exception('Strings must be of equal length.')  
    dist_counter = 0  
    for n in range(len(string1)):  
        if string1[n] != string2[n]:  
            dist_counter += 1  
    return dist_counter / len(string1)
```

The Hamming distance between r1 and q1 is: 0.1690

The Hamming distance between r2 and q1 is: 0.3169

Pairwise Alignment

The Hamming Distance

```
# NCBI Reference Sequence: XP_028905054.1 (platypus hemoglobin subunit A);  
q2 = skbio.Protein("MLTDAEKKEVTALWGKAAGHGEEYGAEALERLFQAFPTTKTYFSHFDLSHGSAQIKAHGKKVADA\  
LSTAAGHFDDMDSALSALSDLHAHKLRVDPVNFKLLAHCILVVLARHCPGEFTPSAHAAMDKFLSKVATVLTISKYR")  
q2
```

Protein

Stats:

length: 141
has gaps: False
has degenerates: False
has definites: True
has stops: False

0 MLTDAEKKEV TALWGKAAGH GEEYGAEALE RLFQAFPTTK TYFSHFDLSH GSAQIKAHGK
60 KVADALSTAA GHFDDMDSAL SALSDLHAHK LRVDPVNFKL LAHCILVVLARHCPGEFTPS
120 AHAAMDKFLS KVATVLTISKY R

```
q2 = skbio.Protein("MLTDAEKKEVTALWGKAAGHGEEYGAEALERLFQAFPTTKTYFSHFDLSHGSAQIKAHGKKVADA\  
LSTAAGHFDDMDSALSALSDLHAHKLRVDPVNFKLLAHCILVVLARHCPGEFTPSAHAAMDKFLSKVATVLTISKYR-")
```


Pairwise Alignment

The Hamming Distance

The Hamming distance between r1 and q2 is: 0.90845

The Hamming distance between r2 and q2 is: 0.92254

```
q2_aligned = skbio.Protein("M-LTDAEKKEVTALWGKAAGHGEEYGAEALERLFQAFPTTKTYFSHFDLSHGSAQIKAHGKKVADA\  
LSTAAGHFDDMDSALSALSDLHAHKLRVDPVNFKLLAHCILVVLARHCPGEFTPSAHAAMDKFLSKVATVLTSKYR")  
print(r1)  
print(q2_aligned)
```

```
MVLSPADKTNVKAAWGKVGAGHAGEYGAEALERMFLSFPTTKTYFPHFDLSHGSAQVKGHGKKVADALTNAVAHVDDMPNALSALSDLHAHKLRVDPV  
M-LTDAEKKEVTALWGKAAGHGEEYGAEALERLFQAFPTTKTYFSHFDLSHGSAQIKAHGKKVADALSTAAGHFDDMDSALSALSDLHAHKLRVDPV
```

The Hamming distance between r1 and q2_aligned is: 0.27465

The Hamming distance between r2 and q2_aligned is: 0.34507


What Is Sequence Alignment?

Mutations:

Substitutions, where one DNA base is replaced with another

Insertions, where one or more contiguous DNA bases are inserted into a sequence

Deletions, where one or more contiguous DNA bases are deleted from a sequence.



CCAAGTT

Simple Align

ACCATGTTT

CCAAGTT

	A	C	C	A	T	G	T	T	T
C		1	1						
C									
A									
A									
G									
T									
T									

Simple Align

ACCATGTTT

CCAAGTT

	A	C	C	A	T	G	T	T	T
C		1	1						
C		1	1						
A									
A									
G									
T									
T									

Simple Align

ACCATGTTT

CCAAGTT

	A	C	C	A	T	G	T	T	T
C		1	1						
C		1	1						
A	1			1					
A									
G									
T									
T									

Simple Align

ACCATGTTT


CCAAGTT

	A	C	C	A	T	G	T	T	T
C		1	1						
C		1	1						
A	1			1					
A	1			1					
G						1			
T					1		1	1	1
T					1		1	1	1

Simple Align

ACCATGTTT

CCAAGTT




ACCATGTTT

-CCAAGTT-

Simple Align

ACCATGTTT

CCAAGTT



ACCA- -TGTTT

-CCAAG- - -TT

Simple Align

ACCATGTTT

CCAAGTT

ACCATGTTT

-CCAAGTT -

$$S = -1+1+1-1+1+1+1-1=4$$


ACCA- -TGTTT

-CCAAG- - -TT


$$S = -1+1+1+1-1-1-1-1-1+1+1=-1$$

Simple Align

Too Simplistic



The Sodium-Potassium Exchange Pump




Simple Align

Too Simplistic

BLOSUM 62 scoring matrix

(positive values are shaded)

A	4																			
R	-1	5																		
N	-2	0	6																	
D	-2	-2	1	6																
C	0	-3	-3	-3	9															
Q	-1	1	0	0	-3	5														
E	-1	0	0	2	-4	2	5													
G	0	-2	0	-1	-3	-2	-2	6												
H	-2	0	1	-1	-3	0	0	-2	8											
I	-1	-3	-3	-3	-1	-3	-3	-4	-3	4										
L	-1	-2	-3	-4	-1	-2	-3	-4	-3	2	4									
K	-1	2	0	-1	-3	1	1	-2	-1	-3	-2	5								
M	-1	-1	-2	-3	-1	0	-2	-3	-2	1	2	-1	5							
F	-2	-3	-3	-3	-2	-3	-3	-3	-1	0	0	-3	0	6						
P	-1	-2	-2	-1	-3	-1	-1	-2	-2	-3	-3	-1	-2	-4	7					
S	1	-1	1	0	-1	0	0	0	-1	-2	-2	0	-1	-2	-1	4				
T	0	-1	0	-1	-1	-1	-1	-2	-2	-1	-1	-1	-1	-2	-1	1	5			
W	-3	-3	-4	-4	-2	-2	-3	-2	-2	-3	-2	-3	-1	1	-4	-3	-2	11		
Y	-2	-2	-2	-3	-2	-1	-2	-3	2	-1	-1	-2	-1	3	-3	-2	-2	2	7	
V	0	-3	-3	-3	-1	-2	-2	-3	-3	3	1	-2	1	-1	-2	-2	0	-3	-1	4
	A	R	N	D	C	Q	E	G	H	I	L	K	M	F	P	S	T	W	Y	V




The values for amino acid substitutions were obtained from Henikoff S & Henikoff JG (1992) Amino acid substitutions matrices from protein blocks. *Proc. Natl. Acad. Sci.* **89**: 10915-10919.

Simple Align

		Seond letter					
		U	C	A	G		
First letter	U	UUU] Phe UUC] UUA] Leu UUG]	UCU] Ser UCC] UCA] UCG]	UAU] Tyr UAC] UAA Stop UAG Stop	UGU] Cys UGC] UGA Stop UGG Trp	U C A G	Third letter
	C	CUU] CUC] Leu CUA] CUG]	CCU] CCC] Pro CCA] CCG]	CAU] His CAC] CAA] Gin CAG]	CGU] CGC] Arg CGA] CGG]	U C A G	
	A	AUU] AUC] Ile AUA] AUG Met	ACU] ACC] Thr ACA] ACG]	AAU] Asn AAC] AAA] Lys AAG]	AGU] Ser AGC] AGA] Arg AGG]	U C A G	
	G	GUU] GUC] Val GUA] GUG]	GCU] GCC] Ala GCA] GCG]	GAU] Asp GAC] GAA] Glu GAG]	GGU] GGC] Gly GGA] GGG]	U C A G	

Simple Align



PAWHEAE

T - The Traceback Matrix

[illegible]

Needleman-Wunsch Alignment


max

$$F(i-1, j-1) + s(c_i, c_j)$$

$$F(i-1, j) - d$$

$$F(i, j-1) - d$$

		H	E	A	G	A	W	G	H	E	E
0	0	-8	-16	-24	-32	-40	-48	-56	-64	-72	-80
P	-8	-2	-9	-17	-25	-33	-41	-49	-57	-65	-73
A	-16	-10	-3	-4	-12	-20	-28	-36	-44	-52	-60
W	-24	-18	-11	-6	-7	-15	-5	-13	-21	-29	-37
H	-32	-14	-18	-13	-8	-9	-13	-7	-3	-11	-19
E	-40	-22	-8	-16	-16	-9	-12	-15	-7	3	-5
A	-48	-30	-16	-3	-11	-11	-12	-12	-15	-5	2
E	-56	-38	-24	-11	-6	-12	-14	-15	-12	-9	1



		H	E	A	G	A	W	G	H	E	E
	•	←	←	←	←	←	←	←	←	←	←
P	↑	↖	↖	←	←	←	←	←	←	←	←
A	↑	↖	↖	↖	←	←	←	←	←	←	←
W	↑	↑	↑	↖	↖	←	↖	←	←	←	←
H	↑	↖	↖	↖	↖	↖	↑	↖	↖	←	←
E	↑	↑	↖	←	↖	↖	↖	↑	↖	↖	←
A	↑	↑	↑	↖	←	↖	↖	↖	↑	↑	↖
E	↑	↑	↖	↑	↖	↖	↖	↖	↖	↖	↖

Needleman-Wunsch Alignment

	H	E	A	G	A	W	G	H	E	E	
0	0	-8	-16	-24	-32	-40	-48	-56	-64	-72	-80
P	-8	-2	-9	-17	-25	-33	-41	-49	-57	-65	-73
A	-16	-10	-3	-4	-12	-20	-28	-36	-44	-52	-60
W	-24	-18	-11	-6	-7	-15	-5	-13	-21	-29	-37
H	-32	-14	-18	-13	-8	-9	-13	-7	-3	-11	-19
E	-40	-22	-8	-16	-16	-9	-12	-15	-7	3	-5
A	-48	-30	-16	-3	-11	-11	-12	-12	-15	-5	2
E	-56	-38	-24	-11	-6	-12	-14	-15	-12	-9	1

	H	E	A	G	A	W	G	H	E	E
P	↑	↖	↖	←	←	←	←	←	←	←
A	↑	↖	↖	↖	←	←	←	←	←	←
W	↑	↑	↑	↖	↖	←	↖	←	←	←
H	↑	↖	↖	↖	↖	↑	↖	↖	←	←
E	↑	↑	↖	←	↖	↖	↖	↑	↖	↖
A	↑	↑	↑	↖	←	↖	↖	↖	↑	↖
E	↑	↑	↖	↑	↖	↖	↖	↖	↖	↖

HEAGAWGHE - E
- PA - -W - HEAE
Score = 1

Local vs. Global Alignment

Global alignment - try to match entire sequences

Useful for closely-related sequences of similar size

Local alignment - allow partial matching

Useful for sequences expected to contain some similarity regions

Global Alignment

Target Sequence
5' ACTACTAGATTACTTACGGATCAGGTACTTTAGAGGCTTGCAACCA 3'
||||| ||||| ||||| ||||| |||||
Query Sequence
5' ACTACTAGATT----ACGGATC--GTACTTTAGAGGCTAGCAACCA 3'

Local Alignment

Target Sequence
5' ACTACTAGATTACTTACGGATCAGGTACTTTAGAGGCTTGCAACCA 3'
|||| ||||| ||||| ||||| |||||
Query Sequence
5' TACTCACGGATGAGGTACTTTAGAGGC 3'

Smith-Waterman Local Alignment

HEAGAWGHEE

PAWHEAE

[illegible]

$$F(0, j) = 0$$

$$T(0, j) = .$$

[illegible]

Smith-Waterman Local Alignment

$$\max \begin{cases} 0 \\ F(i-1, j-1) + s(c_i, c_j) \\ F(i-1, j) - d \\ F(i, j-1) - d \end{cases}$$

AWGHE
AW - HE
Score = 28




		H	E	A	G	A	W	G	H	E	E
0	0	0	0	0	0	0	0	0	0	0	0
P	0	0	0	0	0	0	0	0	0	0	0
A	0	0	0	5	0	5	0	0	0	0	0
W	0	0	0	0	2	0	20	12	4	0	0
H	0	10	2	0	0	0	12	18	22	14	6
E	0	2	16	8	0	0	4	10	18	28	20
A	0	0	8	21	13	5	0	4	10	20	27
E	0	0	6	13	18	12	4	0	4	16	26

		H	E	A	G	A	W	G	H	E	E
.
P
A	.	.	.	↖	.	↖
W	↖	.	↖	←	←	.	.
H	.	↖	←	.	.	.	↑	↖	↖	←	←
E	.	↑	↖	←	.	.	↑	↑	↖	↖	←
A	.	.	↑	↖	←	←	.	↖	↑	↑	↖
E	.	.	↖	↑	↖	↖	←	.	↖	↖	↖

Search

Imagine we have a big book...



... and we want to search it for a specific sentence

It would be
“ so nice if
something
made sense
for a
change.

Lewis Carroll
Alice in Wonderland

Search

- How can we do it in a timely manner?
 - Brute force
 - Indexing
- Do we allow slight changes?
e.g. : *“it **could** be so nice if something made sense”*
- Do we allow insertions and deletions?
e.g. : *“it would be ~~so~~ nice if something made **a little** sense”*
- What if the sentence is repeated in several places in the book?

It would be
“ so nice if
something
made sense
for a
change.
Lewis Carroll
Alice in Wonderland

Sequence Mapping Challenges

Large DBs - millions to billions of nucleotides/AAs

Repetition - biological sequences tend to repeat

Noisy - sequencing errors and real biological variants

BLAST - Basic Local Alignment Search Tool

The most popular alignment tool

BLAST finds regions of similarity between biological sequences.

Compares nucleotide or protein sequences to sequence databases

Calculates the statistical significance of DB hits

Allows searching for **imperfect** sequence matches

Uses a **heuristic** algorithm to improve efficiency



BLAST - Algorithm

1. Index the DB
2. Generate query words
3. compute neighbour words
4. Search the DB for exact word matches - seeds
5. Elongate and combine seeds to get final alignment
6. Score alignment

BLAST - Indexing

Only needed the first time a DB is used

Mask repetitive and low-complexity regions -

ATATATTTATT → atatatttatt

Break DB sequences into overlapping words of length W

- $W=3$ for amino acids
- $W=11$ for nucleotides

Create a lookup table of words with their positions

W T D F G Y P A I L K G G T A C




WTD	1
TDF	2
...	
TAC	14

BLAST - Breaking Query to Words

A query of length L produces $L-W+1$ **overlapping** words of length W

$L = 11$


$W = 3$



BLAST - Finding Neighbour Words

1. For each word, find all neighbourhood words
= words with one change
2. Use a scoring matrix to assign each neighbourhood word a score
3. Discard neighbourhood words with score $< T$


BLAST - Breaking Query to Words



BLAST - Finding Alignment Seeds in DB

Look for **exact** matches of query words with the DB words

Masked regions are ignored





BLAST - Seed Elongation

Elongate each seed to both directions until a score drop $> X$ is encountered

Query:
YANCL**EH**KMG S

$$X = 2$$



BLAST - Scoring the Alignment

Calculate total alignment score

A	N	C	Q	E	H	K	M	G
A	P	C	Q	E	H	K	R	G
4	-2	9	5	5	8	5	-1	6

Discard alignments with score $< S$

Remaining alignments are called High scoring Sequence Pairs - **HSPs**

BLAST - Scoring the Alignment

- Calculate alignment **bit score**
 - Independent of query length
 - Independent of DB size

$$S' = \frac{\lambda S - \ln(K)}{\ln(2)}$$

- Calculate **E-value** - the number of hits with score $\geq s$ that one can expect to find in DB by chance

$$E = \frac{L \times N}{2^{S'}}$$

L - query length , N - DB length , S' - bit score

Smaller $E \rightarrow$ better hit

BLAST - Scoring the Alignment

W – word size (query and DB)

T – neighborhood words score cutoff

X – allowed score drop during seed elongation

S – HSP score cutoff



What would happen if
we **increase** T ?



Scale and Speed

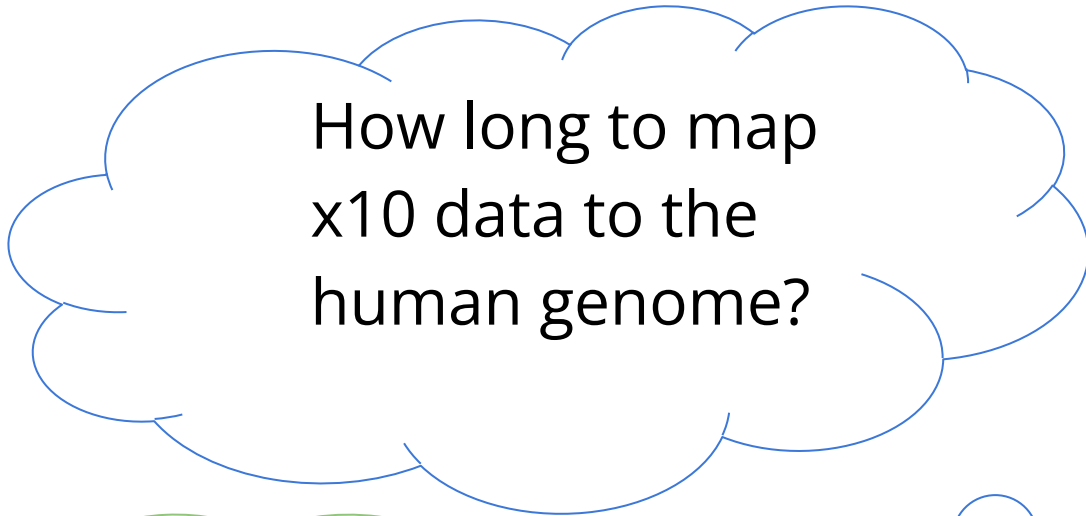
We need to map millions to hundreds of millions of reads

Can we use Blast?

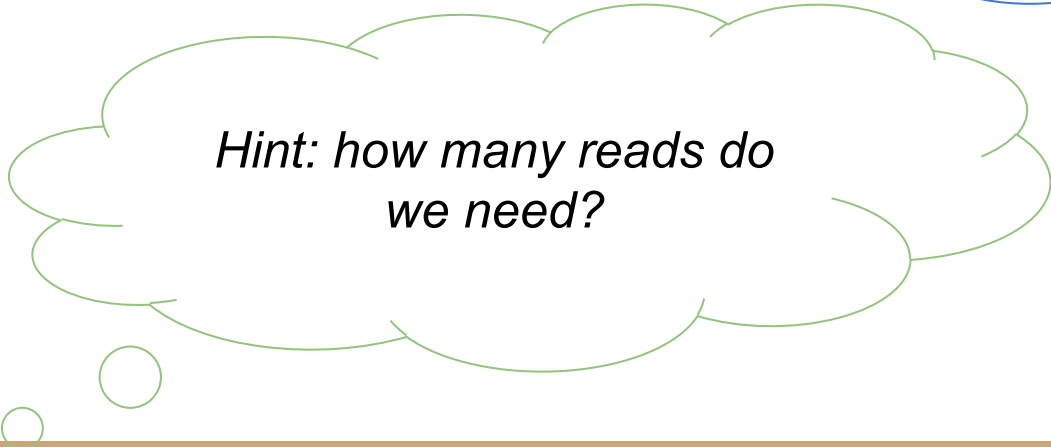
Blastn - ~100 reads / sec

Human genome - ~ 3Gb

Assume 100bp reads



How long to map
x10 data to the
human genome?



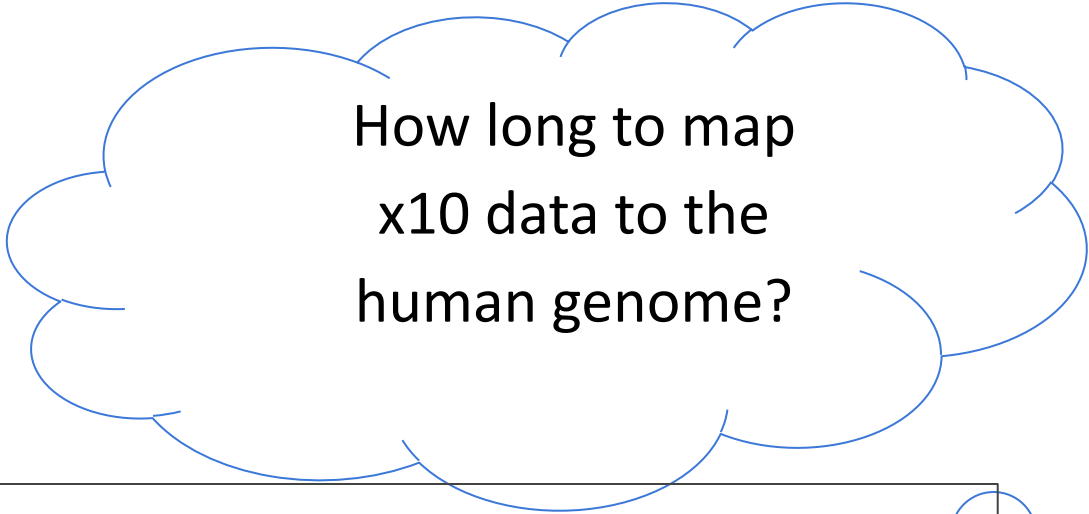
*Hint: how many reads do
we need?*

Can We Use Blast in NGS?

Blastn - ~100 reads / sec

Human genome - ~ 3Gb

Assume 100bp reads



How long to map
x10 data to the
human genome?

Data required:

3 Gb x 10 = 30 Gb

Reads required:

30 Gb / 100 = 300 M reads

Time to map:

300 M reads / (100 reads/sec) = 3M sec = ~ **35 days**

BWA - Burrows-Wheeler Aligner

Specifically designed for mapping of short reads

Maps ~2,200 reads / sec (one CPU)

Allows parallel computing

Contains three algorithms - the most useful is **BWA-MEM**

BWA - Limitations

Only works for nucleotides (usually DNA, not RNA)

Less effective when:

- Queries are very long
- Reads are highly diverged from the reference
- Reads contain lots of sequencing errors

Usually offers a good accuracy-speed balance

BWA - Algorithm Overview

Step 1: Index the reference genome


Step 2: Search for reads

Indexing is based on the **Burrows-Wheeler's transformation**

Index allows easy searching:

- Quick
- Memory efficient

BWA - The Burrows Wheeler Transform



BWT(abracadabra\$) = ard\$rcaaaabb

BWA - The Burrows Wheeler Transform

BWT is **reversible** - we can get back from BWT(G) to G

BWT(G) tends to cluster the same characters together - easy to compress

BWT(abracadabra\$) = ard\$rcaaaabb

Using some additional data structures, BWT(G) can be searched efficiently

BWA - The Burrows Wheeler Transform

1. Create index of reference genome:

Input: reference in fasta format

```
$ bwa index genome.fasta
```

2. Map reads to reference:

Input: reads file or pair (for PE data) in fastq format

```
$ bwa mem genome.fasta reads_R1.fq reads_R2.fq -o aln.sam
```